

1. AWS Shared Responsibility Model

- **AWS Responsibility:**

Physical Security: Protection of AWS data centers, servers, and physical infrastructure.

Network and Hardware Security: Ensuring secure network and hardware infrastructure.

Infrastructure Management: Management of cloud infrastructure (e.g., computing, storage, networking).

Compliance: Ensuring the AWS infrastructure complies with regulations and standards (e.g., SOC, ISO).

- **Customer Responsibility:** Securing data (encryption, backup, etc.).

Data Protection: Managing IAM roles, policies, and permissions.

Access Management: Securing operating systems on services like EC2.

Operating System Security: Managing and securing applications running on AWS services.

Application Security: Ensuring the customer's environment complies with applicable regulations (e.g., GDPR, HIPAA).

Compliance: ensuring compliance in the cloud adherence to security, privacy, and regulatory standards relevant to their industry or region.

2. EC2(Amazon Elastic Compute Cloud) Instance Types

Instance Type	Cost	Use Case	Flexibility	Reliability
On-Demand	Pay-as-you-go (higher cost)	Applications with unpredictable traffic	High (no commitment)	Moderate: No guarantee of availability; can be terminated with short notice during high demand.
Spot	Significant discount (variable)	Fault-tolerant or flexible workloads	Limited (may be interrupted)	Low: Instances can be interrupted with two-minute warning if AWS needs the capacity. Suitable only for fault-tolerant workloads.
Reserved	Significant discount (1-3 years)	Steady-state workloads with predictable demand	Moderate (commitment required)	High: Reserved instances are guaranteed for the term commitment, providing stable capacity and availability.
Dedicated Hosts	Pay per host (higher cost)	Compliance-driven workloads with specific licensing	Low (dedicated physical server)	High: Full control over physical server, ensuring no resource sharing with other AWS customers, leading to predictable performance.
Savings Plans	Flexible long-term savings	Workloads that require flexibility across services	High (flexible across services and regions)	High: Provides consistent capacity over the term commitment, but with more flexibility than Reserved Instances.
Dedicated Instances	Pay for dedicated physical server	Compliance needs requiring physical isolation	Moderate (dedicated hardware but shared instances)	Moderate to High: Offers some isolation from other tenants but still shares the physical server hardware. Predictability is better than Spot or On-Demand, but not as high as Dedicated Hosts.

- **On-Demand** and **Spot Instances** offer flexibility but have **moderate to low reliability** due to potential interruptions (Spot) or lack of guarantees (On-Demand).
- **Reserved Instances** and **Savings Plans** provide **high reliability** due to long-term commitments, though Savings Plans offer more flexibility in instance family, region, and service.
- **Dedicated Hosts** offer **the highest reliability** with dedicated physical servers, suitable for workloads requiring strict compliance and isolation.
- **Dedicated Instances** provide a middle ground, offering a reasonable balance of isolation and reliability, but at a higher cost than shared instances.

3. EC2 Family

Category	Instance Family	Primary Use Case
General Purpose	M Series, T Series	Balanced compute, memory, and networking for web apps, dev/test, and small databases.
Compute Optimized	C Series	CPU-bound applications, such as high-performance web servers and scientific computing.
Memory Optimized	R Series, X Series	Memory-intensive applications like large databases and in-memory analytics.
Storage Optimized	I Series, D Series, H Series	Storage-heavy workloads like NoSQL databases, big data, and data warehouses.

Accelerated Computing	P Series, G Series, F Series	GPU-based or FPGA-accelerated tasks like machine learning, graphics rendering, and hardware acceleration.
------------------------------	------------------------------	---

4. EC2 Placement Group Types

Placement Group Type	Purpose	Network Latency	Fault Isolation	Max Instances Per AZ
Cluster	Low-latency, high-throughput communication	Low	No isolation from hardware failure	Limited by available resources
Spread	High availability, isolated from hardware failure	Moderate	High (spread across hardware)	7 instances per AZ
Partition	Fault isolation, used for large-scale apps	Moderate to High	High (isolated partitions)	Limited by partition size

- **Cluster Placement:** High-performance applications that need fast, high-throughput communication, like scientific simulations.
- **Spread Placement:** Applications that need to be resilient to hardware failures, such as fault-tolerant systems.
- **Partition Placement:** Large-scale distributed applications that require isolated fault domains, like large Hadoop clusters.

5. Amazon S3(**Simple Storage Service**) **Object Lock** is a feature designed to protect objects from being accidentally or intentionally **deleted** or **altered**. It enforces a *write-once-read-many* (WORM) model, ensuring data immutability for a specified retention period.

6. Amazon S3 Server-Side Encryption (SSE) Types

Feature	SSE-S3 (S3 Managed Keys)	SSE-KMS (AWS KMS Managed Keys)	SSE-C (Customer-Provided Keys)
Key Management	Managed automatically by Amazon S3	Managed through AWS KMS; users define policies	Managed by the user; key must be provided with each request
Security Level	Basic encryption	Enhanced security with audit logging and key rotation	Maximum control (user fully responsible for key management)
Key Control	Not user-controllable	User-defined key policies and permissions	Fully user-controlled
Operational Complexity	Lowest (fully managed)	Moderate (requires KMS and IAM setup)	Highest (users must supply and secure keys)
Cost	No additional cost	Additional cost for KMS key usage and requests	No AWS cost, but key management overhead for the user
Compliance Support	Basic compliance	Meets stringent compliance standards (e.g., PCI DSS, HIPAA)	Suited for custom compliance or self-managed key needs
Advantages	Simple, cost-effective	Advanced security and auditing	Complete control and flexibility
Use Cases	General-purpose encryption	Scenarios requiring stringent access control or auditing	Scenarios requiring full user control over keys

- **SSE-S3:** Ideal for low-cost, straightforward encryption needs where simplicity is a priority.
- **SSE-KMS:** Suitable for scenarios requiring enhanced security, access control, and compliance standards.
- **SSE-C:** Best for organizations that demand complete control over encryption keys, often for custom compliance or unique security needs.

7. Amazon S3 Storage Classes

Storage Class	Description	Use Cases	Storage Cost	Retrieval Latency	Minimum Storage Duration
S3 Standard	Designed for data that is frequently accessed and active.	Active data for websites, analytics, or mobile apps	High	Milliseconds	None
S3 Standard-IA	Cost-effective storage for infrequently accessed data with rapid access requirements.	Backup or archival data that requires fast access	Medium	Milliseconds	30 days

S3 Glacier Instant Retrieval	Low-cost storage for rarely accessed cold data, with rapid retrieval.	Archive data that may need fast access occasionally	Lower	Milliseconds	90 days
S3 Glacier Flexible Retrieval	Cost-effective storage for extremely rarely accessed data, with flexible retrieval.	Archive (e.g., compliance data, historical records)	Even lower	Minutes to hours	90 days
S3 Glacier Deep Archive	Lowest-cost option for data that is rarely accessed and must be retained long-term.	Long-term archiving (e.g., contracts, photos)	Lowest	12–48 hours	180 days

- **S3 Standard:** Real-time transactional data, frequently accessed customer files, or media content.
- **S3 Standard-IA:** Monthly backup data or operational logs accessed occasionally.
- **S3 Glacier Instant Retrieval:** Fast-access archive data such as scanned medical records.
- **S3 Glacier Flexible Retrieval:** Legal records or regulatory files that require occasional retrieval.
- **S3 Glacier Deep Archive:** Long-term storage of compliance documents or historical media.

8. **S3 Bucket names** must be **globally unique** across all AWS accounts in all regions.

9. AWS Storage Services

Category	Object Storage	Block Storage	File Storage	
Service	Amazon S3	Amazon EBS	Amazon EFS	Amazon FSx
Purpose	Store unstructured data with lifecycle management, versioning, and cross-region replication.	Persistent storage for EC2 instances, supports snapshots and encryption with performance tiers.	Distributed file system supporting NFS protocol, allows simultaneous access by multiple EC2 instances .	High-performance file systems (Windows File Server, Lustre) for enterprise needs.
Typical Use Cases	Static website hosting , backups, data lakes, content distribution.	Databases, transactional applications, file systems.	Big data analytics, CMS, multi-instance shared files.	Enterprise file sharing, HPC .
Access Method	REST API , SDKs, AWS CLI, Management Console.	EC2 instance attachment via API , CLI, Management Console.	NFS mount targets via EC2 or on-premises with Direct Connect or VPN.	SMB (FSx for Windows), Lustre APIs, or Linux NFS.
High Availability (HA)	Highly durable (11 9s), supports cross-region replication for HA.	Single AZ , but snapshots provide cross-AZ backup and restore options.	Multi-AZ by default (One Zone for cost-saving options).	Multi-AZ by default, depending on configuration.

10. Private Subnet Devices or Services in a VPC Access the Internet

- Using a NAT Gateway (Recommended)
NAT Gateway is a managed service that deployed in the public subnet to provide internet access for devices in the private subnet. The private subnet's route table points to the NAT Gateway as the default route for outbound traffic.
- Using a NAT Instance (Deprecated)
An EC2 instance in the public subnet serves as a NAT device.
NAT instances require manual configuration for availability and scaling.
You can configure port forwarding on a NAT instance by modifying its security group and adding port forwarding rules on the instance itself (e.g., using iptables or other firewall configuration tools).
- Using a Proxy Server
Devices in the private subnet access the internet through a proxy server in the public subnet.
Requires additional configuration for traffic management.
- Using VPC Endpoints (AWS Services Only)

Private subnet devices access AWS services (e.g., S3, DynamoDB) through VPC Endpoints without internet exposure.

This applies only to specific AWS services.

- In contrast, **NAT Gateways** and **NAT Instances** are used for **IPv4** traffic from private subnets. The **Egress-Only Internet Gateway** is the corresponding resource when dealing with **IPv6**.

11. Stateful Firewall and Stateless Firewall

- **Stateful Firewall**: Security Groups (automatically allow return traffic for established connections, deny any default).
- **Stateless Firewall**: NACLs (require explicit rules for both inbound and outbound traffic, deny any default).

12. Steps to Enable Internet Access for EC2 in a Private Subnet

- Create a NAT Gateway or NAT Instance in the Public Subnet
- Update the Private Subnet's Route Table
- Ensure the Public Subnet's Route Table Has a Route to the Internet Gateway
- Check Security Groups and Network ACLs
- IAM Role Permissions (Optional)

13. In AWS, **private IP addresses** assigned to resources (such as EC2 instances) typically remain **persistent** during the **lifetime** of the resource.

14. Use **identity-based policies** to control permissions at the **user** or **role level** across resources.

15. Use **resource-based policies** to enforce access rules **directly on resources**, especially for cross-account or fine-grained access control.

16. AWS Organizations is a service that allows central management of multiple **AWS accounts** within a single organization. It provides tools for account management, billing, and policy control.

17. To **encrypt** all **communications** from **clients** through **ELB** to **web servers**, you can:

- Deploy server certificates on both **ELB** and **instances** for end-to-end **encryption** with TLS termination at the ELB.
- Deploy server certificates **only** on the **instances** and configure **ELB** for **TCP** forwarding to achieve end-to-end encryption without TLS termination at the ELB.

18. Elastic Load Balancers (ELB)

- Automatically monitors the health of registered targets and **routes** traffic **only** to **healthy targets**.
- ELB spans **multiple Availability Zones** (AZs) by default, ensuring fault tolerance and seamless failover.

19. Auto Scaling Strategies

Scaling Strategy	Description
Target Tracking Scaling	Automatically adjusts the number of instances to maintain a specific metric (e.g., CPU utilization).
Step Scaling	Defines actions based on the severity of metric changes , with different scaling amounts for different thresholds.
Simple Scaling	Triggers scaling actions based on a single threshold crossing, scaling out or in by a set number of instances .
Scheduled Scaling	Allows you to scale based on a schedule , useful for predictable increases or decreases in demand (e.g., daily or weekly traffic patterns).

20. AWS Auto Scaling combined with Lifecycle Hooks

Enables you to perform custom actions during instance scaling events, such as when EC2 instances are launching or terminating. The lifecycle hook gives you control over the scaling process, allowing you to pause an instance as it is being launched or terminated, and perform tasks like configuration, validation, or graceful shutdown.

Functionality	Description	Use Case
Custom initialization during launch	When an instance is launched as part of an Auto Scaling group, you can pause the instance at the "Pending" state to complete custom setup actions (e.g., configuring software, running tests).	Allows you to ensure that all necessary configuration is done before the instance begins accepting traffic.

Graceful termination	When an instance is terminating, you can pause it at the “Terminating” state to perform clean-up tasks (e.g., saving logs, deregistering from load balancers, draining connections).	Ensures that the instance is properly removed from service and that no work is lost during shutdown.
Pause scaling operations for validation	Before scaling out or in, lifecycle hooks can be used to pause operations to ensure that the environment is stable, and that additional validation or checks are performed.	Ensures that scaling operations are only completed after critical steps, such as database backups or state checks, are validated.

21. **SNI (Server Name Indication)** is an extension of the TLS (Transport Layer Security) protocol that allows a client (such as a web browser) to indicate the hostname it is attempting to connect to during the TLS handshake process. This enables a **single server** (or load balancer, like ELB) to **host multiple SSL/TLS certificates** for **different domains** on the **same IP address** and **port**.

22. The Serverless AWS Services

Category		Service	Description
Compute	Serverless	AWS Lambda	Run code in response to events; pay only for usage.
		AWS Fargate	Serverless compute for containerized applications.
	Server-based	Amazon EC2	Virtual servers in the cloud with full control over operating systems and configurations.
		Amazon LightSail	Simple, cost-effective servers for smaller workloads or projects.
Storage	Serverless	Amazon S3	Object storage with scalability and pay-as-you-go pricing.
		Amazon EFS (with ECS/Fargate)	Serverless file storage for applications.
	Server-based	Amazon EBS	Persistent block storage for EC2 instances.
		Amazon FSx	Fully managed file systems (e.g., Windows File Server, Lustre).
		Amazon S3 Glacier	Long-term, low-cost data archiving solution.
Databases	Serverless	Amazon DynamoDB	NoSQL database service with global replication and on-demand scalability.
		Amazon Aurora Serverless	Auto-scaling relational database service.
	Server-based	Amazon RDS	Managed relational database service supporting multiple database engines.
		Amazon Aurora	High-performance relational database compatible with MySQL and PostgreSQL.
		Amazon Redshift	Data warehousing solution optimized for large-scale analytics.
Networking	Serverless	Amazon API Gateway	Create and manage APIs at scale.
		AWS App Runner	Deploy containerized web applications or APIs.
	Server-based	Amazon VPC	Provides a private network environment for server-based resources.
		Elastic Load Balancing (ELB)	Distributes traffic across EC2 instances, containers, and Lambda.
		AWS Direct Connect	Dedicated network connection to AWS for high bandwidth.
		Amazon Route 53	DNS web service for domain routing.
Security	Serverless	AWS Secrets Manager	Securely manage and rotate sensitive information.
	Server-based	AWS IAM	Manages user access and permissions to AWS resources.
		AWS Shield	Protects against DDoS attacks.
		AWS WAF	Web Application Firewall for controlling incoming traffic to web applications.
Deployment	Serverless	AWS Step Functions	Orchestrate serverless workflows.
AWS CodePipeline		Automate CI/CD pipelines for applications, including serverless ones.	
Messaging		Amazon SQS	Fully managed message queues (Standard & FIFO).
		Amazon SNS	Managed pub/sub messaging for notifications and alerts.

		Amazon EventBridge	Serverless event bus for application integration.
Analytics		AWS Glue	Serverless data integration and ETL service.
		Amazon Athena	Query data in S3 using SQL.
		Amazon Kinesis	Serverless data streaming for real-time analytics.
Machine Learning		Amazon Rekognition	Image and video analysis service.
		Amazon Lex	Build chatbots and conversational interfaces.
Monitoring & Management	Server-based	Amazon CloudWatch	Monitors resources and applications with metrics, logs, and alarms.
		AWS Config	Tracks changes in resources and ensures compliance.
AWS Systems Manager		Centralized management for AWS resources and operational tasks.	
Application Services		Elastic Beanstalk	Simplifies application deployment and management on AWS.
		Amazon MQ	Managed message broker for Apache ActiveMQ.

23. Comparison of (AWS Lambda with S3) and (ECS Fargate with EFS)

Feature	AWS Lambda + S3	ECS Fargate + EFS
Use Case	Lightweight, event-driven, serverless tasks.	Persistent workloads requiring high availability and file sharing.
Storage Type	Amazon S3 for storage, accessed via API (no native mount).	Amazon EFS, POSIX-compliant file system for temporary/shared access.
Storage Capacity	Limited by Lambda's 10 GB /tmp storage, insufficient.	EFS provides scalable storage, meeting the 50 GB+ requirement.
Performance	Higher latency as S3 is accessed over the network.	Low latency, high performance for frequent file operations.
Operational Complexity	Lambda auto-scales but lacks large storage integration.	Fargate simplifies cluster management, supporting direct EFS mounts.
Scalability	Auto-scaling Lambda but limited by its storage constraints.	Auto-scaling Fargate tasks with robust storage capabilities via EFS.
Cost	Lower cost but S3 frequent access may increase expenses.	Higher cost but better value for high I/O workloads .

24. ECS Three Key Components

Component	Detailed Description	Purpose
Task Definition	A JSON file that acts as a blueprint for your containerized application. - It specifies: - Docker image : The container image to use. - Resource requirements : CPU and memory allocation per container. - Networking : Port mappings and protocols. - Environment variables : Application-specific settings. - Logging : Integration with services like Amazon CloudWatch. - Volume definitions : Persistent storage settings for containers.	Ensures consistent deployment configurations across multiple environments or clusters.
Clusters	A logical grouping of compute resources managed by ECS. - Resources include: - EC2 instances : When using EC2 launch type for hosting containers. - AWS Fargate : For serverless container hosting. - A cluster can span across multiple Availability Zones (AZs). - Cluster management is automated, including scaling and workload distribution.	Centralized organization of compute resources for running containerized applications while ensuring fault tolerance.
Services	A service is responsible for maintaining the desired number of tasks from a task definition. - Features include: - Auto Scaling : Dynamically adjusts the number of running tasks based on demand. - Service Discovery : Enables automatic discovery of containerized services. - Load Balancing : Integrates with ALB or NLB to route traffic to running tasks. - Ensures rolling updates and health checks for tasks. - Supports zero-downtime deployments.	Maintains high availability, balances traffic and provides automated recovery for containerized workloads.

25. Enhance the confidentiality of data stored in Amazon EBS

- **Enable EBS Encryption:**

Encrypt data at rest, in transit, and during snapshots. Use Amazon-managed keys (**AWS KMS**) or customer-managed keys for encryption.

- **Snapshot Encryption:**

Encrypted snapshots inherit encryption from the source volume, ensuring encrypted backup data.

26. Four Categories of Amazon FSx

Amazon FSx Service	Supported Protocols	Key Features	Ideal Use Cases
FSx for Windows File Server	SMB	Fully managed Windows file system, Active Directory integration, Windows ACLs, file-level access auditing.	Windows-based applications, shared file storage, enterprise workloads.
FSx for Lustre	Lustre	High-performance, scalable, and optimized for compute-intensive workloads, supports parallel file access.	Machine learning, big data analytics, high-performance computing (HPC).
FSx for NetApp ONTAP	NFS, SMB, iSCSI	Fully managed NetApp ONTAP system, data protection, encryption, multi-protocol support, cloud data management.	Hybrid cloud, enterprise applications requiring high availability and scalability.
FSx for OpenZFS	NFS, SMB, iSCSI	Managed OpenZFS system with features like compression, snapshots, data cloning, and fast recovery.	Developers needing advanced data protection, high performance, and ZFS compatibility.

27. AWS Storage Gateway is a hybrid cloud storage service that allows you to seamlessly connect your on-premises environment with cloud storage. It is designed to integrate on-premises IT environments with Amazon Web Services (AWS) to provide a secure and scalable way to store and access data in the cloud.

Type of Volume Gateway	Description	Use Case
Cached Volumes	Data is stored primarily in Amazon S3 , and frequently used data is cached locally on the gateway.	Ideal for workloads with large datasets that require frequent access to a subset of data.
Stored Volumes	Entire dataset is stored locally on-premises, with backups of the data asynchronously replicated to Amazon S3.	Best for workloads that need low-latency access to all data and require reliable off-site backups.

28. Comparison of Databases

Service	Type	Key Features	Use Cases	Scalability	Capacity
RDS	Relational DB	Fully managed, supports MySQL, PostgreSQL, Oracle, SQL Server, MariaDB	OLTP, transactional apps, ERP systems	manual scaling , Vertical scaling with instance resizing	64 TB (MySQL, PostgreSQL), 16 TB (SQL Server, Oracle)
Aurora		Fully managed, Compatible with MySQL/PostgreSQL, high availability, fault-tolerant	Mission-critical apps needing high performance	auto-scaling , Scales reads with replicas	Up to 128 TB (for Aurora)
Redshift	Data warehouse (SQL)	Columnar storage, OLAP optimization, Redshift Spectrum for S3 queries, encryption by default	BI, data analytics , complex reporting	Scales with compute nodes, only one leader node	Up to petabytes of data (scalable, MPP architecture)
DynamoDB	NoSQL (key-value)	Fully managed, Serverless, low-latency , built-in backup, on-demand scaling, encryption by default	IoT , gaming leaderboards, mobile apps	Auto-scaling, Horizontal (unlimited write/read scaling)	Virtually unlimited (auto-scaling)
ElastiCache	In-memory	Redis/Memcached support, sub-millisecond latency, session caching	Real-time analytics, gaming leaderboards	Horizontal scaling with shards/replicas	

Neptune	Graph DB	Supports RDF/SPARQL, property graphs, fully managed	Social networks, fraud detection, knowledge graphs	Horizontal (sharded replication)	
Timestream	Time-series DB	Optimized for time-series data, serverless scaling, tiered storage	IoT, telemetry, industrial sensors	Scales writes and storage automatically	

- **Performance:** Aurora and DynamoDB are ideal for high-performance workloads.
- **Scalability:** DynamoDB and Timestream offer unlimited horizontal scalability.
- **Specific Use Cases:** Choose Neptune for graphs, Timestream for time-series, and Redshift for analytics.
- Require the Use of an **Encrypted Replica** to enable encryption: RDS, Aurora, RedShift, DynamoDB

29. Amazon Aurora Architecture and Its Features

- **Architecture:**
 - **Storage Layer:**
Aurora uses a **distributed storage architecture** that replicates data across six copies in three Availability Zones (AZs). The storage layer is **log-structured** and optimized for database workloads, providing durability and high availability. Storage automatically scales in 10 GB increments up to 128 TB without requiring manual intervention.
 - **Compute Layer:**
The compute instances handle query execution and business logic. Aurora supports **reader and writer nodes**:
The **writer node** handles write requests and updates the shared storage.
Reader nodes handle read-only requests, with automatic failover if the writer fails.
 - **Connection Layer:**
Aurora provides a **cluster endpoint** for connecting applications to the appropriate database instance, ensuring seamless load balancing and high availability.
- **Key Features:**
 - **Performance and Scalability:**
Aurora delivers up to **5x the performance** of standard MySQL and 3x that of PostgreSQL. Storage and compute scale independently, accommodating varying workload demands.
 - **Fault Tolerance and High Availability:**
Data is replicated across multiple AZs, and the service automatically recovers from hardware failures or storage issues. It supports **automatic failover** within 30 seconds, ensuring minimal downtime.
 - **Backup and Durability:**
Continuous backups to **Amazon S3** with point-in-time recovery. No performance degradation during backups.
 - **Read Scalability:**
Aurora supports up to 15 **read replicas**, which can scale read-heavy workloads while maintaining low-latency performance.
 - **Cost Efficiency:**
Aurora is designed to deliver enterprise-grade performance and features at a **fraction of the cost** of commercial databases like Oracle or SQL Server.
 - **Security:**
Offers encryption for data at rest and in transit using **AWS Key Management Service (KMS)**. Integration with **Amazon VPC** for network isolation and **IAM** for access control.

- **Integration with AWS Services:**

Aurora integrates seamlessly with AWS services like **AWS Lambda**, **Amazon CloudWatch**, and **AWS Glue**, allowing advanced analytics, monitoring, and automation.

30. Amazon Redshift Architecture

Feature	Leader Node	Compute Node
Function	Coordinates query execution and management.	Stores data, executes queries, and performs calculations.
Data Storage	Does not store data.	Stores data in slices, with each node storing a part of the data.
Primary Task	Manages query parsing, optimization, and distribution.	Executes data scans, joins, filters, and processes computations.
Query Handling	Receives queries, parses them, generates execution plans, and distributes them to Compute Nodes.	Processes assigned query tasks and returns results to the Leader Node.
Parallel Processing	Does not participate in parallel data processing.	Supports parallel query processing by working on different slices.
Resources	No storage or computation resources.	Has its own CPU, memory, and storage resources for computation and data handling.
Role in Redshift	Central coordinator that oversees query execution.	Performs the actual computation and storage tasks.

31. Amazon DynamoDB architecture

- **Distributed Architecture**

DynamoDB operates on a distributed architecture, meaning the data is not stored on a single server or node. Instead, the data is spread across multiple physical nodes that work together to provide high availability and horizontal scalability.

- **Tables and Partitions**

Tables: In DynamoDB, data is stored in tables, and each table contains multiple items. Each table has a primary key that uniquely identifies each item.

Partitions: DynamoDB automatically splits and distributes data across multiple partitions. Each partition stores a subset of the data. As the size of a table grows, DynamoDB automatically adds more partitions to handle the increased data and requests.

- **Primary Key**

Simple Primary Key: Composed of a single attribute (partition key) used to uniquely identify an item in the table.

Composite Primary Key: Consists of two attributes, with the first being the partition key and the second being the sort of key. This allows multiple items in the same partition to be ordered by the sort of key.

- **Read and Write Capacity**

Read Capacity Units (RCUs) and Write Capacity Units (WCUs): DynamoDB manages throughput using provisioned capacity or on-demand capacity modes.

On-Demand Mode: DynamoDB automatically adjusts throughput based on actual usage, suitable for unpredictable workloads.

Provisioned Mode: Users define the required throughput at the time of table creation, suitable for stable workloads.

- **Global Distribution (Global Tables)**

DynamoDB supports **Global Tables**, which replicate data automatically across multiple AWS regions, ensuring low latency and high availability in cross-region applications.

- **Auto Scaling**

DynamoDB features **Auto Scaling**, which automatically adjusts read and write capacity based on workload demand, helping to handle traffic spikes without manual intervention.

- **Transactional Support**

DynamoDB offers **ACID** (Atomicity, Consistency, Isolation, Durability) transaction support, enabling users to perform transactional operations across multiple items, ensuring data consistency.

- **Consistency Models**

DynamoDB provides two consistency options:

Strongly Consistent Reads: Ensures that the read returns the most recent write.

Eventually Consistent Reads: Allows reads that may not reflect the most recent write but provides lower latency and higher throughput.

- **Indexes**

DynamoDB supports several types of indexes for query optimization:

Global Secondary Index (GSI): Allows queries based on attributes other than the primary key.

Local Secondary Index (LSI): Enables different sorting of data within the same partition key, optimizing queries on the same partition.

- **Data Encryption**

DynamoDB has built-in data encryption, and all data stored in DynamoDB can be automatically encrypted without user intervention. It uses AWS Key Management Service (KMS) for key management.

- **DynamoDB Streams** captures changes to data in DynamoDB tables and provides a time-ordered sequence of these changes. Essentially, DynamoDB Streams allows you to track modifications (such as inserts, updates, and deletes) made to items in a DynamoDB table. You can then process and react to these changes programmatically, enabling event-driven architectures and other use cases that require real-time data processing.

32. ElastiCache Database Engines:

Feature	Redis	Memcached
Data Structures	Advanced (e.g., lists, sets, hashes, sorted sets)	Simple key-value store
Persistence	Yes (AOF and RDB persistence)	No
Clustering	Yes (Redis Cluster)	Yes (Auto-Sharding)
Encryption	Yes (in transit and at rest)	No (only at transit with external solutions)
Pub/Sub Messaging	Yes	No
Use Cases	Caching, real-time analytics, session management, queues	Caching, session management, and temporary storage